
MODULE 2 – WORD AND TEXT REPRESENTATIONS

1. Word Relationships

Natural Language relies heavily on relationships between words.

Understanding these helps NLP systems find **semantic similarity** (words that mean alike) and **contextual meaning** (words that often appear together).

◆ A. What Are Word Relationships?

Words are not isolated — their meaning depends on their relation to other words.

Examples:

- “doctor” relates to “hospital”, “patient”.
 - “car” and “automobile” are similar in meaning.
 - “hot” is opposite of “cold”.
-

◆ B. Types of Word Relationships

Type	Description	Example
Synonymy	Words with similar meanings.	<i>begin–start, big–large</i>
Antonymy	Opposite meanings.	<i>good–bad, fast–slow</i>
Polysemy	Same word, multiple meanings.	“bank” → river bank / financial bank
Homonymy	Same spelling/sound, unrelated meanings.	“bat” → animal / cricket bat
Hyponymy	“Is-a” relationship (subcategory).	<i>rose is a flower</i>
Hypernymy	General category relationship.	<i>fruit → apple, mango</i>
Collocation	Words that often occur together.	“make decision”, “strong tea”

◆ C. The Distributional Hypothesis

“Words that occur in the same contexts tend to have similar meanings.”

Example:

- “dog” and “cat” appear in similar contexts like *pet, feed, animal*.
- Thus, they should be represented closely in vector space.

This forms the **basis for word embeddings** (Word2Vec, GloVe, etc.).

💡 2. Word Embeddings

◆ A. What is a Word Embedding?

A **word embedding** represents words as **dense vectors of real numbers** (e.g., 50–300 dimensions).

These vectors are arranged so that **semantically similar words are close** in the vector space.

Word Simplified 3D Vector

King [0.8, 0.6, 0.1]

Queen [0.7, 0.7, 0.2]

Man [0.9, 0.5, 0.0]

Woman [0.8, 0.6, 0.1]

👉 Relationship:

[

King - Man + Woman ≈ Queen

]

This is known as **vector arithmetic**, showing that embeddings capture **meaning and analogy**.

◆ B. Why Word Embeddings?

Earlier methods like **Bag of Words (BoW)** or **TF-IDF** treat words as independent tokens:

- No context.
- No semantic relationship.
- Very high dimensional.

Word embeddings solve these issues by:

- Representing meaning geometrically.
 - Being compact and efficient.
 - Capturing relationships automatically.
-

◆ C. Comparison: Bag of Words vs Word Embedding

Aspect	Bag of Words	Word Embedding
Representation	Count-based	Dense numerical vectors
Context	Ignored	Preserved
Dimension	High (1 per word)	Low (50–300)
Semantics	None	Captured
Example	“dog” ≠ “cat”	“dog” ≈ “cat”

✳️ 3. Word Embedding Techniques

✿ A. Bag of Words (BoW)

Concept:

Represents a document by the **frequency of each word**, ignoring grammar and order.

Steps:

1. Build a vocabulary of unique words.
2. Count occurrences of each word in every document.

Example:

Sentence	Words
----------	-------

S1: “I love NLP” I, love, NLP

S2: “I love AI” I, love, AI

Vocabulary: [I, love, NLP, AI]

Document 1 love NLP AI

S1 1 1 1 0

S2 1 1 0 1

Pros:

- Simple to understand
- Works well for small datasets

Cons:

- Ignores order and context
 - Sparse (large, mostly zero vectors)
-

B. TF-IDF (Term Frequency – Inverse Document Frequency)

Improves BoW by weighting words according to their **importance**.

◆ Formula

```
[  
TF(t, d) = \frac{\text{count of term t in document d}}{\text{total words in document d}}  
]  
[  
IDF(t) = \log\left(\frac{N}{\text{number of documents containing t}}\right)  
]  
[  
TF-IDF = TF × IDF  
]
```

◆ Example

Word Document Frequency IDF (\log_{10})

the	100	0
-----	-----	---

machine	5	1.3
---------	---	-----

learning	3	1.7
----------	---	-----

If “learning” appears 4 times in one document:

```
[  
TF-IDF(learning) = 4 × 1.7 = 6.8  
]
```

Meaning:

Rare, topic-specific words get higher weight.

Pros:

- Highlights meaningful terms
- Reduces weight of common words

Cons:

- Ignores word order and semantics
 - High-dimensional sparse matrix
-

C. Word2Vec

Developed by **Mikolov et al., 2013 (Google)**

Uses **neural networks** to learn word vectors that encode meaning.

◆ **Two Architectures**

1. **CBOW (Continuous Bag of Words)**

- Predicts target word from surrounding context.
- Example: context = “I __ NLP” → predict “love”.

2. **Skip-Gram**

- Predicts surrounding context from target word.
 - Example: target = “NLP” → predict (“I”, “love”).
-

◆ **How It Works**

Each word is represented by a **one-hot vector** (like [0,0,1,0,0,...]).

The model learns **hidden layer weights** that become dense word vectors.

◆ **Example in Python (Gensim)**

```
from gensim.models import Word2Vec
```

```
sentences = [
```

```
    ['machine', 'learning', 'is', 'fun'],
```

```
[['python', 'makes', 'machine', 'learning', 'easy'],  
 ['artificial', 'intelligence', 'includes', 'machine', 'learning']]
```

```
model = Word2Vec(sentences, vector_size=50, window=3, min_count=1, sg=1)
```

```
print(model.wv['machine']) # vector representation
```

```
print(model.wv.most_similar('learning'))
```

Output (example):

```
Vector for 'machine': [0.01, -0.02, 0.03, ...]
```

```
Most similar to 'learning': [('python', 0.13), ('intelligence', 0.12)]
```

◆ Optimization Techniques in Word2Vec

1. Negative Sampling

- Updates only a few weights per training sample.
- Faster training, works well for large vocabularies.

2. Hierarchical Softmax

- Uses a binary tree to compute probabilities efficiently.

3. Window Size:

- Controls how many surrounding words to consider as context (e.g., 2–5 typical).

4. Training Method:

- CBOW is faster; Skip-Gram works better for rare words.
-

◆ Advantages

- ✓ Captures semantic and syntactic relations
- ✓ Compact representation
- ✓ Supports analogy reasoning (King–Man+Woman=Queen)

◆ Limitations

- ✖ Needs large text corpus
 - ✖ Context window is fixed (doesn't capture long-range meaning)
-

D. Other Embedding Models

Model	Description	Example Use
GloVe (Global Vectors)	Combines global word co-occurrence with local context.	Semantic analysis
FastText	Represents words as subword (character n-gram) vectors → handles misspellings or rare words.	“play”, “player”, “playing” share features

4. Simple N-Gram Models

◆ A. Definition

An **N-gram** is a sequence of *N consecutive words* used to model text.

N Example (Sentence: “I love NLP”)

1-gram I, love, NLP

2-gram I love, love NLP

3-gram I love NLP

◆ B. How N-Gram Models Work

They predict the probability of a word given the previous (N-1) words.

[
P(w_n | w_{n-1}, ..., w_{n-N+1})
]

Example (Bigram Model):

[

```
P(\text{NLP} | \text{love}) = \frac{\text{Count}(love, NLP)}{\text{Count}(love)}
```

If:

- Count(love NLP) = 5
- Count(love) = 20

Then:

```
[  
P(\text{NLP} | \text{love}) = 5 / 20 = 0.25  
]
```

◆ C. Applications

- Text prediction and autocomplete
 - Speech recognition
 - Statistical machine translation
 - Spelling correction
-

◆ D. Limitations

- High memory usage for large N.
 - Data sparsity (some N-grams may never appear).
 - Cannot handle long-range dependencies.
-



5. Comparison of All Techniques

Method	Type	Context Captured	Meaning Captured	Dimensionality
BoW	Count-based	✗ No	✗ No	High
TF-IDF	Weighted count	✗ No	⚠ Partial	High
Word2Vec	Neural	✓ Yes	✓ Yes	Low
N-Gram	Statistical	✓ Local	⚠ Partial	Variable

6. Summary

Concept	Core Idea	Example
Word Relationships	Synonymy, antonymy, and semantic similarity between words.	car–automobile
BoW	Count frequency of words.	[I:1, love:1, NLP:1]
TF-IDF	Weighted frequency to highlight unique words.	“learning”: high weight
Word2Vec	Neural model capturing meaning in dense vectors.	King– Man+Woman=Queen
Optimization	Negative sampling, hierarchical softmax.	Efficient training
N-Gram	Predict next word using context.	“I love → NLP”

In Summary:

Word and text representation methods convert language into numbers so that algorithms can analyze meaning and context.

From simple counts (**BoW**, **TF-IDF**) to neural embeddings (**Word2Vec**, **GloVe**), and statistical models (**N-gram**), these form the foundation for all modern NLP tasks.